

Smart Traffic Management Database

Chrysikos Christos, 9432 | Mavridis Antonis, 9563 | Gitopoulos Giorgos, 9344

Databases @ ECE AUTh, November 2021

1. Introduction

This project deals with the theoretical design of a database of a proposed application. After the description of the application and its data requirements, we will present the user types and the structure of the database (*Entity - Relationship model* and *Relational model*). Moreover, some examples will be used in order to explain the database functionality.

1.1. Application Description

The proposed application concerns *traffic management* in a *smart city*. A *smart city* is a technologically modern urban area that uses different types of electronic methods, voice activation methods and sensors to collect specific data. That data is used to improve the operations across the city, such as traffic management.

The purpose of the application is to match each vehicle with the fastest feasible route in order to reach the desirable destination and help the drivers to locate available parking slots easier. Also, it aims to detect traffic violations and store them. First of all, we assume a computing cloud that will perform all the necessary computations for the application. In addition, a database will be needed to store all the collected data.

Concerning the application, the exact position of each vehicle will be collected by GPS, it will be stored in the database and a traffic metric will be computed for each area of the city according to the number of vehicles in the area. Every driver will select its destination and the cloud will export the optimal route for this vehicle, using a *Shortest Path Algorithm* (i.e. *Dijkstra's Algorithm*) and taking into account the neighboring areas, the traffic metric of each area, the traffic lights and the tolls. Also, vision sensors will be placed in every parking slot and the data of their status will be stored to the database. So, the drivers will be able to find the empty parking slots in a specific area using the application. About the traffic violations, cameras will be placed through the city streets and in the traffic lights too, in order to detect speed limit or road signs violations and violations of red traffic lights respectively.

1.2. Data Requirements

Some of the data that the database should store, with the corresponding estimations of storage size per day, are:

- the users with their id, name, age, gender and mobile number (500,000 ~ 10,000,000 users),
- the vehicles with their licence plate, type, driver id, location, destination and route (~ 1 vehicle per user),
- the areas of the city with their id, name and traffic metric (20 ~ 50 areas),
- the couples of neighboring areas in order to declare the routes (~ 3 neighboring areas per area),
- the routes that are exported by the cloud (as a sequence of areas *) with their id and duration (~ 3 routes per vehicle),
- the traffic lights with their details (5,000 ~ 20,000 traffic lights),
- the tolls with their details (3 ~ 10 tolls),
- the parking slots with their details (1,000,000 ~ 20,000,000 parking slots)
- the traffic violations and their details (5,000 ~ 20,000 violations)

* Note that the smaller each area is, the more precise the route is. For simplicity we assume larger areas of a city.

2. User Categories & Requirements

In this section, we will present the types of users of the database and their corresponding data access requirements.

- **End user:** The end users of the application are the drivers of the vehicles. They can access all the data of their own vehicle, their own route, the data of every area of the city, the tolls, the traffic lights and the data of every parking slot. They are able to write in the database just their own destination, while all the other data are read - only for them.
- **Administrator:** The administrators can read and modify all the data of the database. They are responsible for the proper operation of the system and they should take action if they detect any type of problem.
- **Violation officer:** Police officers responsible for traffic violations will have read and write rights on the traffic violations and read rights on the information of the users who commit the violations.

- **Sensor - Embedded system:** The several sensors with their embedded systems will be considered as users of the database, as they are able to write and modify data. For example, a system of a parking slot sensor updates the status of the slot in the database, every time that it changes and the embedded system of a camera writes down the driver id (automatically found by the plate number of the vehicle) and its fee in the violations table, when it detects that it violated the speed limit. The embedded systems of the sensors have no access to other data and cannot read data at all. They can just write in or modify some specific memory addresses.
- **Cloud:** The computing cloud has access to all the data concerned to the route computation. It can read the current location of every vehicle to compute the traffic metric of every area as a fuzzy value and writes it in the table of the areas. Also, it reads the location of the traffic lights, the tolls, the destination of every vehicle, the neighboring areas and the computed traffic metric of each area in order to compute the fastest route for each vehicle. When it does that, it writes the route id in the corresponding line of the vehicle table, the route id and its duration in the route table and the areas that the route passes through in the route - area table (Note that the tables will be presented in depth in section 4).

3. Entity – Relationship Model

3.1. Description

The **entities** of the **Entity - Relationship** model are: *User*, *Area*, *Route*, *Vehicle*, *Violations*, *Tolls*, *Parking Slot*, and *Traffic Light*. For every *User* there is at least one *Vehicle*, each *Vehicle* can only be located in one *Area* at a time and can follow only one *Route* at time, each *Vehicle* must have a starting and a destination *Area*, each *Route* consists of many different *Areas* and each *Area* can be a part of many different *Routes*. An *Area* can have *Parking Slots*, *Tolls* and *Traffic Lights*. Also, each *Area* has its own neighboring *Areas* and each *User* can commit one or more *Violations*.

The description of the **attributes** of our **entities** follows:

- For every *User* there must be a *name*, a *gender*, a unique *id* which is the primary key, an *age* and a *mobile number*.
- Every *Vehicle* must have a unique *license plate* and a *type* which defines the type of the vehicle (car, motorcycle, truck, etc.).
- Every *Area* has a unique *id* which is the primary key, a *name* and must have a *traffic metric* which shows the traffic situation in the *Area*.
- Every *Route* has a unique *id* which is the primary key and a *duration* which determines how long the *Route* is estimated to last (in minutes).
- Every *Violation* has a unique *id*, a *fee* and a *type* (red light, stop sign, speeding etc...).
- A *Parking Slot* has a unique *id*, an exact *location* and a *status* which shows if it is occupied or not.
- *Traffic Light* has a unique *id*, a *status* (red or green), an exact *location* and a *duration* (in minutes).
- *Tolls* have an *id*, a *location* and a *toll* which defines the price of the *Tolls*.

Some **assumptions** that will be made are:

- Every *Vehicle* has a unique *User*. There can be no *Vehicle* with two different owners.
- A *Vehicle* can have alternating locations depending on time. For example, if a *Vehicle* traverses through a *Route* it will change from two to multiple locations in order to arrive in the final destination.
- An *Area* can have multiple *Vehicles* simultaneously.
- The current location of each *Vehicle* is being updated by a *GPS tracking unit* in a fixed time interval schedule.
- The *Route* of each *Vehicle* is decided by the *computing cloud* that uses a *Shortest Path Algorithm*, taking into account the *traffic metrics*, the *Traffic Lights* and the *Tolls* of the *Areas*. For example, the *User* with *id* = 1, starts from *Area* with *id* = 3, chooses the destination *Area* with *id* = 10, the *computing cloud*, taking into consideration the parameters that were mentioned previously, decides that the optimal *Route* is the one with *id* = 2, which passes through *Areas* with *id* = 1, 4, 6, 10 and each *Area* is neighboring to the previous and the next one.
- Two different *Vehicles* can be matched with the same *Route*.
- The *traffic metric* of an *Area* is computed as a fuzzy value after counting the *Vehicles* that are located in the *Area* and is updated every a standard amount of minutes.
- The *status* of every *Parking Slot* is being updated by the embedded systems of the *sensors* that are placed in the *Parking Slot*.
- The *Violations* are recorded by *cameras* and are automatically written to the *database*.
- The *Neighboring Areas* (see section 4) define all the couple of *Areas*, so that a *Vehicle* can pass from the first to the second and are used by the *computing cloud* to define valid *Routes*. For example, a *Route* can start from *Area* with *id* = 1, pass through *Area* with *id* = 2 and reach *Area* with *id* = 3, only if there are the *Neighboring Areas* couples 1-2 and 2-3 in the respective table.
- If the *cloud* computes that the optimal *Route* for a *Vehicle* is the one mentioned in the previous bullet, it will declare an *id* for that *Route* (i.e. *id* = 105), it will write that *id* in the table of *Vehicles* next to the specific *Vehicle* and it will also add the *id* of the *Route* with all the *ids* of the *Areas* it passes through in a *Route - Areas* table (it will be described in section 4). So, it will write the couples (105, 1), (105, 2), (105, 3).
- As noted in subsection 1.2 too, we consider large *Areas* for simplicity. The optimal approach would be to consider short pieces of streets as *Areas* in order to increase the precision of the *Routes*.

3.2. Entities

In the next tables, our **entities** are explained further:

Entity Name	User	Entity Name	Vehicle
	Entity where we store every user in the Database		Entity where we save every Vehicle in the Database
	Strong Entity		Weak Entity - Cannot exist without a User
	<ul style="list-style-type: none"> • <u>id</u> • name • age • gender • mobile_number 		<ul style="list-style-type: none"> • <u>license_plate</u> • type
Entity Name	Area	Entity Name	Route
	Entity where we store every Area in the Database		Entity where we store every Route in the Database
	Strong Entity		Strong entity
	<ul style="list-style-type: none"> • <u>id</u> • name • traffic_metric 		<ul style="list-style-type: none"> • <u>id</u> • destination
Entity Name	Violation	Entity Name	Parking Slot
	Entity where we save every traffic Violation in the Database		Entity where we store every Parking Slot in the Database
	Weak Entity - Cannot exist without User		Weak Entity - Demands an Area
	<ul style="list-style-type: none"> • <u>id</u> • fee • type 		<ul style="list-style-type: none"> • <u>id</u> • location • status
Entity Name	Traffic Light	Entity Name	Tolls
	Entity where we store every Traffic Light in the Database		Entity where we save every Toll in the Data Base
	Weak Entity - Demands an Area		Weak Entity - Demands an Area
	<ul style="list-style-type: none"> • <u>id</u> • status • location • duration 		<ul style="list-style-type: none"> • <u>id</u> • location • toll

Figure 1: *Entities*

3.3. Relationships

The relationships between the defined *entities* are presented below:

Relationship Name	User - own - Vehicle	Entity Name	Vehicle - follow - Route
Description	Every user must have 1 to many Vehicles and every vehicle must have a unique user as owner	Description	Many Vehicles can share the same Route.
Properties	Owes, Identifying	Properties	Follow, Identifying
Cardinality	1:N	Cardinality	N:1
Dependency	User fully dependant, vehicle partial dependant	Dependency	Vehicle fully dependant, route partial dependant
Attributes	-	Attributes	-

Relationship Name	Vehicle - is located in - Area	Relationship Name	Vehicle - has a starting and a destination - Area
Description	There can be many Vehicles in the same Area	Description	Many different Vehicles can share the same starting and destination Area.
Properties	Located in, Identifying	Properties	has-A, Identifying
Cardinality	N:1	Cardinality	N:1
Dependency	Vehicle fully dependant, Area partial dependant	Dependency	Vehicle fully dependant, Area partial dependant
Attributes	-	Attributes	-

Relationship Name	Routes - pass through - Areas	Relationship Name	Is - neighboring - to
Description	Many Areas can be a part of a Route and many routes can share the same Area	Description	Every Area can be a neighbor to any other Area
Properties	Pass through, Identifying	Properties	Neighboring, Unary
Cardinality	M:N	Cardinality	M:N
Dependency	Route fully dependant, Area partial dependant	Dependency	Area partial dependant
Attributes	-	Attributes	-

Figure 2: Relationships (i)

Relationship Name	Area - has - Parking Slot	Entity Name	Area - has - Tolls
Description	Every Area can have from 0 to many Parking Slots	Description	Every Area can have from 0 to many Tolls
Properties	has-A, Identifying	Properties	has-A, Identifying
Cardinality	1:N	Cardinality	1:N
Dependency	Area partial dependant, Parking Slot fully dependant	Dependency	Area partial dependant, Tolls fully dependant
Attributes	-	Attributes	-

Entity Name	Area - has - Traffic Light	Relationship Name	User - commit - Violation
Description	Every Area can have from 0 to many Traffic Light	Description	Every user can have 0 to many Violations
Properties	has-A, Identifying	Properties	Commits, Identifying
Cardinality	1:N	Cardinality	1:N
Dependency	Area partial dependant, Traffic Light fully dependant	Dependency	Violation fully dependant, User partial dependant
Attributes	-	Attributes	-

Figure 3: Relationships (ii)

3.4. Entity – Relationship Diagram

The **entity - relationship** diagram of the *database* based on the defined *entities* and *relationships*, according to *Chen's notation* is the following:

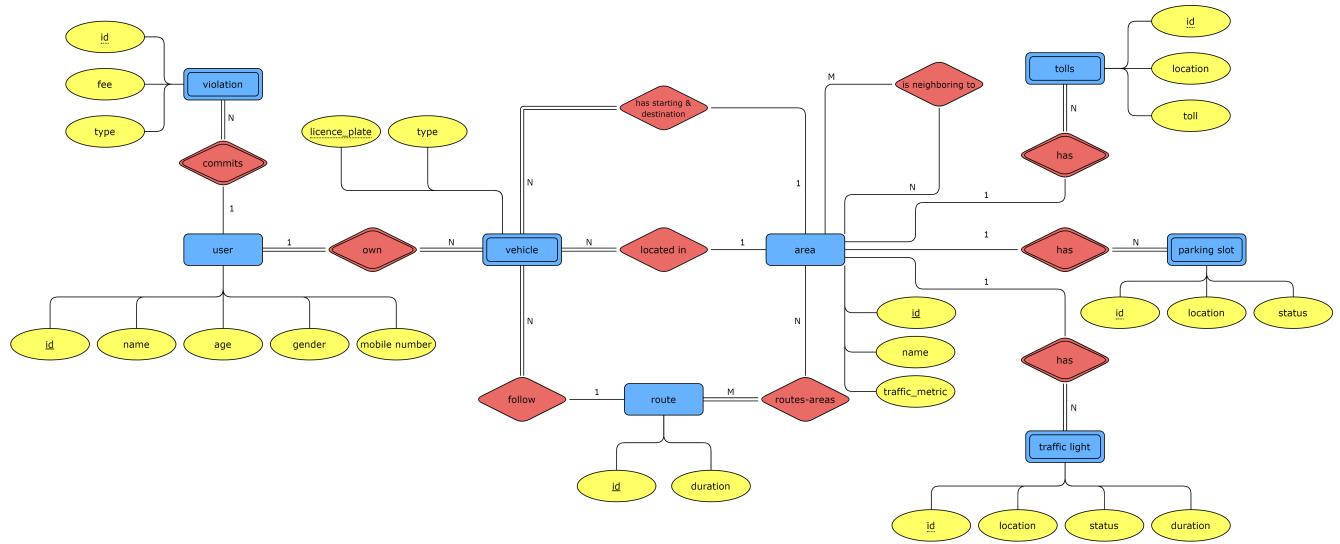


Figure 4: Entity – Relationship Diagram

4. Relational Model

4.1. Data Domains

The **data domains** of the *relational model* are presented below:

Domain	Data Type
id	FIXED_CHAR (10)
name	VARIABLE_CHAR (30)
age	INTEGER_RANGE [18,99]
gender	ENUMERATED {MALE, FEMALE}
mobile_number	VARIABLE_CHAR (10)
integer	INT
violation_type	ENUMERATED {TRAFFIC_LIGHT, SPEED_LIMIT, TRAFFIC_SIGN}
vehicle_type	ENUMERATED {CAR, MOTORCYCLE, TRUCK}
money	FLOAT
metric	ENUMERATED {VERY_LOW, LOW, NORMAL, HIGH, VERY_HIGH}
duration	INT
location	VARIABLE_CHAR (24)
parking_slot_status	ENUMERATED {EMPTY, NOT_EMPTY}
traffic_light_status	ENUMERATED {GREEN, RED}

Figure 5: Data Domains

4.2. Relations

The **relations** of the *relational model* are shown in the next tables:

User		Vehicle	
attributes :		attributes :	
Name	Domain	Name	Domain
id	id	licence_plate	id
name	name	type	vehicle_type
age	age	user_id	id
gender	gender	route_id	id
mobile_number	mobile_number	current_area_id	id
constraints :		constraints :	
Primary Key	id	Primary Key	licence_plate & user_id
Foreign Key	-	Foreign Key	user_id ~> User route_id ~> Route current_area_id ~> Area starting_area_id ~> Area destination_area_id ~> Area

Area		Route	
attributes :		attributes :	
Name	Domain	Name	Domain
id	id	id	id
name	name	duration	time
traffic_metric	metric	constraints :	
Primary Key	id	Primary Key	id
Foreign Key	-	Foreign Key	-

Figure 6: *Relations (i)*

Violation		Traffic Light	
attributes :			
Name	Domain	Name	Domain
id	id	id	id
type	violation_type	status	traffic_light_status
fee	money	location	location
user_id	id	duration	time
constraints :		constraints :	
Primary Key	id & user_id	Primary Key	id & area_id
Foreign Key	user_id ~> User	Foreign Key	area_id ~> Area
Parking Slot		Tolls	
attributes :			
Name	Domain	Name	Domain
id	id	id	id
location	location	location	location
status	parking_slot_status	toll	money
area_id	id	area_id	id
constraints :		constraints :	
Primary Key	id & area_id	Primary Key	id & area_id
Foreign Key	area_id ~> Area	Foreign Key	area_id ~> Area
Routes - Areas		Neighboring Areas	
attributes :			
Name	Domain	Name	Domain
route_id	id	area1_id	id
area_id	id	area2_id	id
constraints :		constraints :	
Primary Key	route_id & area_id	Primary Key	area1_id & area2_id
Foreign Key	route_id ~> Route area_id ~> Area	Foreign Key	area1_id ~> Area area2_id ~> Area

Figure 7: Relations (ii)

4.3. Relational Diagram

The relational diagram follows:

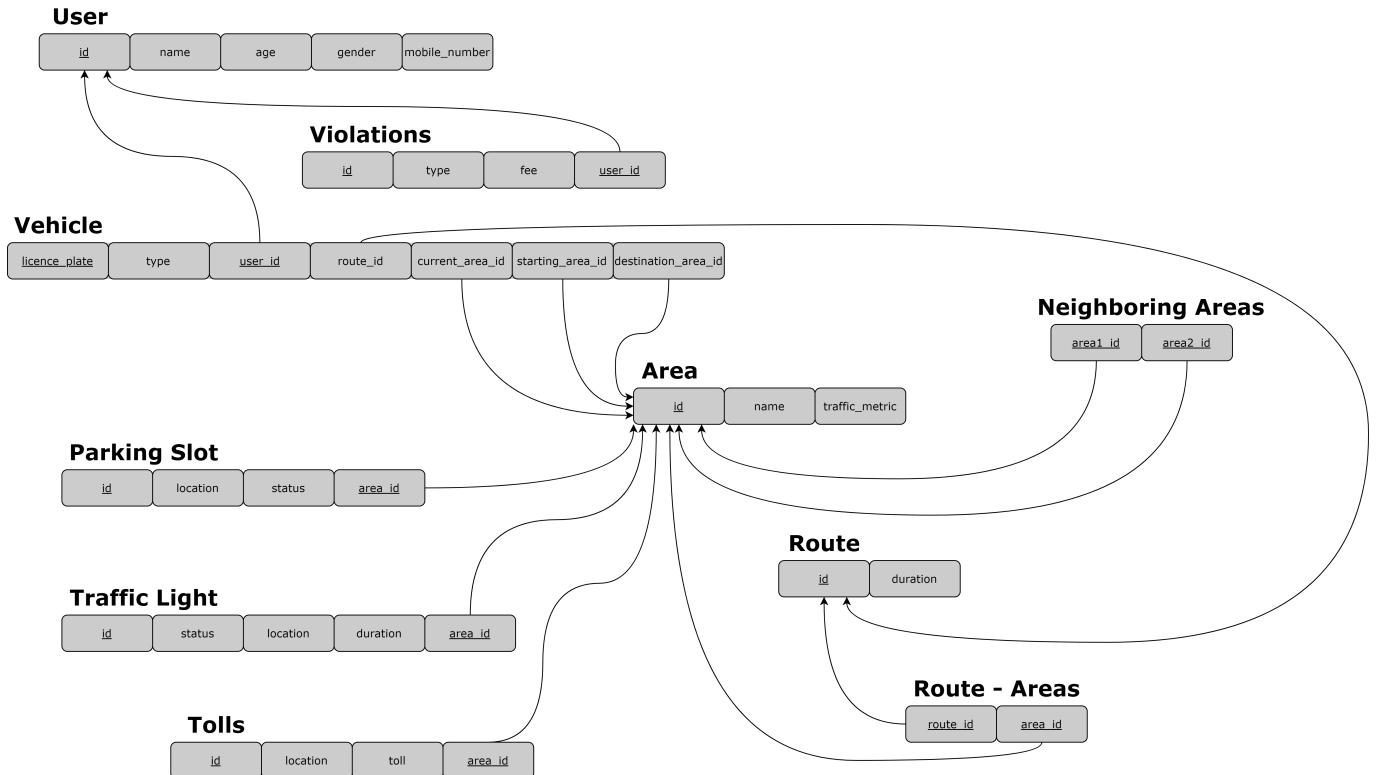


Figure 8: Relational Diagram

4.4. Views

Now, we will define **database views** in order to store some data that will be accessed by the *database users* frequently. The first view will consist of all the *Areas names* and the empty *Parking Slots* in that *Area*. It will be used by the *Users* who are trying to locate an empty *Parking Slot* in a specific *Area*. The view is expressed in *relational algebra* as:

$$\left[\pi_{name,id}(Area) \right] \bowtie \left[\pi_{area_id,location} (\sigma_{status=EMPTY} (ParkingSlot)) \right]$$

or

$$\pi_{name,area_id,location} \left[\sigma_{id=area_id} \left[[\pi_{name,id}(Area)] \times [\pi_{area_id,location} (\sigma_{status=EMPTY} (ParkingSlot))] \right] \right]$$

Another view to be defined, is the *Users* who commit *Violations* with the *Violation details*. It will be accessed by *Violation Officers* in order to see the information of *Users* who committed every *Violation*. In *relational algebra* it is defined as:

$$\pi_{name,age,gender,mobile_number,type,fee,user_id} \left[\sigma_{id=user_id} [(User) \times \pi_{type,fee,user_id}(Violation)] \right]$$

5. Examples

5.1. Table Examples

Subsequently, we will present an example of an **instance** of the *database*, that is all the data stored in the tables at a particular moment of time. The estimation for the stored data size can be found in 1.2.

User				
id	name	age	gendre	mobile_number
AI05069901	Jim Carrey	32	MALE	6972461286
IK13114254	Steve Smith	26	MALE	6921673161
KE01019933	Chloe Brown	36	FEMALE	6981620227
BN08028872	John Davis	52	MALE	6915623763
OM27126296	Marie Miller	21	FEMALE	6993357512

Vehicle						
licence_plate	type	user_id	route_id	current_area_id	starting_area_id	destination_area_id
GRTH639633	CAR	AI05069901	THPAFA3163	THTO920993	THPA237310	THFA021224
GRTH699821	CAR	IK13114254	THPOTO0912	THFA021224	THPO651209	THTO920993
GRTH098331	MOTORCYCLE	KE01019933	THKETO3441	THFA021224	THKE835612	THTO920993
GRTH707342	TRUCK	BN08028872	THPOPA0121	THKE835612	THPO651209	THPA237310
GRTH788242	CAR	OM27126296	THKEPO9884	THKE835612	THKE835612	THPO651209

Violation			
id	type	fee	user_id
PP3781AI01	TRAFFIC_SIGN	20	KE01019933
SA9312IK54	TRAFFIC_LIGHT	300	IK13114254
KF9571KE33	SPEED_LIMIT	90	KE01019933
YT1267BN72	SPEED_LIMIT	60	OM27126296
EZ5205OM96	TRAFFIC_SIGN	40	OM27126296

Area		
id	name	traffic_metric
THKE835612	Center	VERY_HIGH
THPA237310	Panorama	LOW
THTO920993	Toumpa	HIGH
THFA021224	Faliro	NORMAL
THPO651209	Polichni	VERY_LOW

Route	
id	duration (min)
THPAFA3163	20
THPOTO0912	25
THKETO3441	15
THPOPA0121	30
THKEPO9884	15

Figure 9: Database Instance (i)

Parking Slot			
id	location	status	area_id
THKE412320	41.427462,23.652820	EMPTY	THPA237310
THPA292425	39.523312,24.535225	NOT_EMPTY	THPA237310
THTO421920	42.327402,19.865920	NOT_EMPTY	THTO920993
THFA402225	40.623412,22.955825	EMPTY	THFA021224
THPO282225	38.628712,22.005225	NOT_EMPTY	THTO920993

Traffic Light				
id	status	location	duration (sec)	area_id
THKETL3267	GREEN	40.424462,22.654820	35	THKE835612
THPATL2572	RED	40.887462,21.692820	25	THPA237310
THTOTL1903	RED	39.426662,22.677820	20	THTO920993
THFATL0212	RED	42.427462,24.652880	5	THFA021224
THPOTL9032	GREEN	40.400062,23.699920	30	THPO651209

Tolls			
id	location	toll	area_id
THKET95127	41.488462,21.659920	1.2	THKE835612
THPAT09978	40.517462,21.882820	0.7	THPA237310
THTOT12842	39.462662,22.677820	1	THTO920993
THFAT09051	42.429962,24.611880	1.1	THFA021224
THPOT65481	41.473062,23.699920	0.9	THPO651209

Neighboring Areas	
area1_id	area2_id
THPO651209	THAM581274
THKE835612	THSY681273
THKE835612	THFA021224
THTO920993	THFA021224
THKE835612	THPO651209

Route - Areas	
route_id	area_id
THKETO3441	THFA021224
THKETO3441	THTO920993
THPOPA0121	THKE835612
THPOPA0121	THTO920993
THPOPA0121	THPA237310

Figure 10: Database Instance (ii)

5.2. Relational Algebra Examples